

CSC 450 – COMPUTER NETWORKS

Lecture 3

Design Philosophy and Applications

HTTP – Hyper Text Transfer Protocol

Recap

- Layering
- Protocol Stacks
- OSI 7 layer protocol
 - We also know that not all 7 layers are applied
- Important Protocols (overview)
 - IP
 - TCP
 - UDP
- Brief overview of Internet
 - History
 - Services
 - Reliability measures

Today's Class

- A brief overview of Applications
 - Application Programs
 - Application Protocols
- A detailed description of the Application Protocol HTTP

History of the Web as an Application

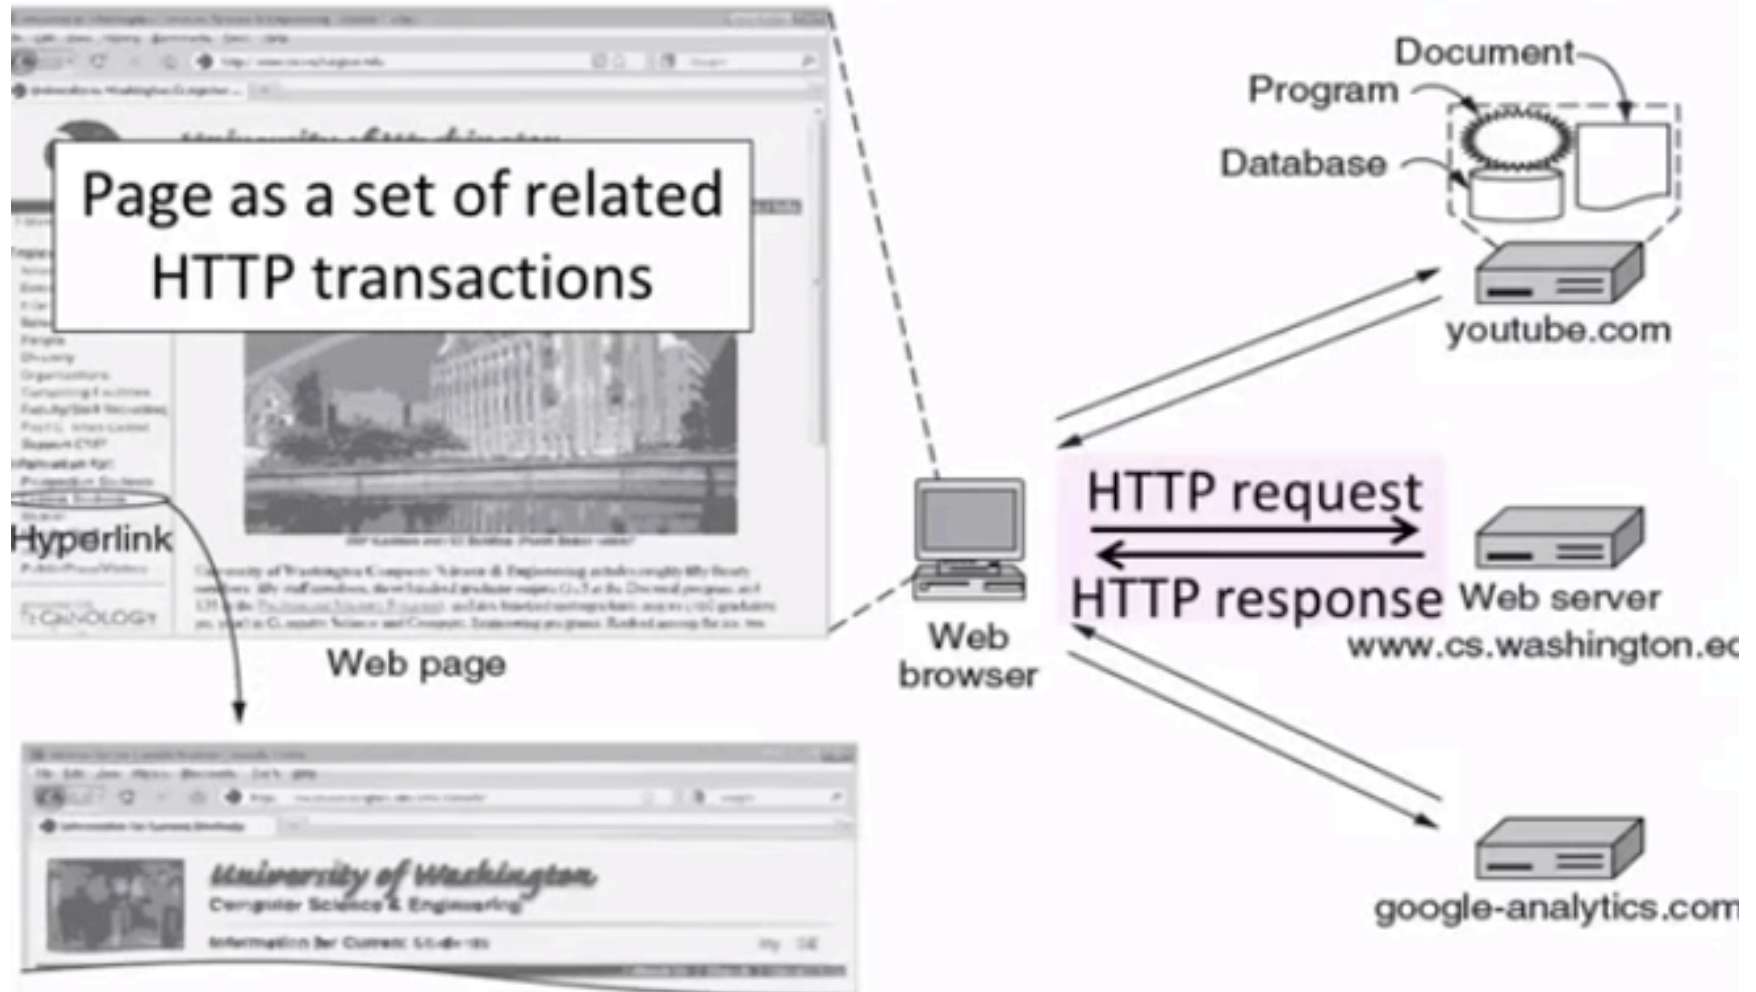
- Sir Tim Berners-Lee (1955-)

- Inventor of the Web
 - Dominant Internet app since mid 90s
 - He now directs the W3C
- Developed Web at CERN in '89
 - Browser, server and first HTTP
 - Popularized via Mosaic ('93), Netscape
 - First WWW conference in '94 ...




Source: By Paul Clarke, CC-BY-2.0, via Wikimedia Commons

Web Context



HTTP Context

- HTTP is a request/response protocol for fetching Web resources
 - Runs on TCP, typically port 80
 - Part of browser/server app



Fetching a Web page with HTTP

- Start with the page URL:

`http://en.wikipedia.org/wiki/Vegemite`

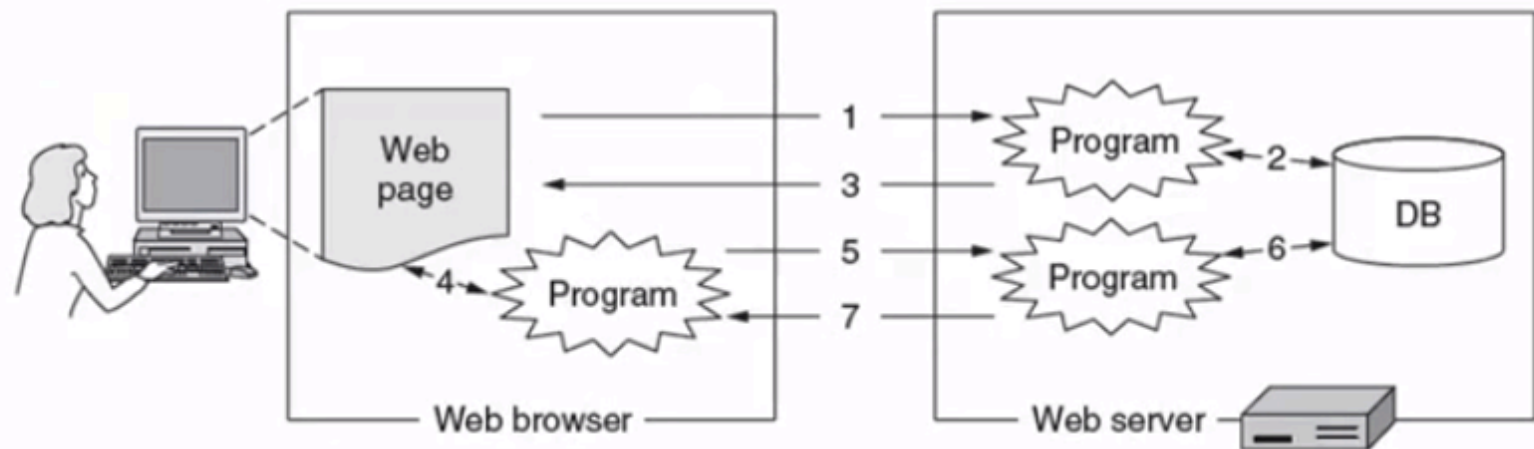
The URL `http://en.wikipedia.org/wiki/Vegemite` is broken down into three parts by purple brackets below it:

- Protocol:** `http`
- Server:** `://en.wikipedia.org`
- Page on server:** `/wiki/Vegemite`

- Steps:
 - Resolve the server to IP address (DNS)
 - Set up TCP connection to the server
 - Send HTTP request for the page
 - (Await HTTP response for the page)
 - * * Execute / fetch embedded resources / render
 - Clean up any idle TCP connections

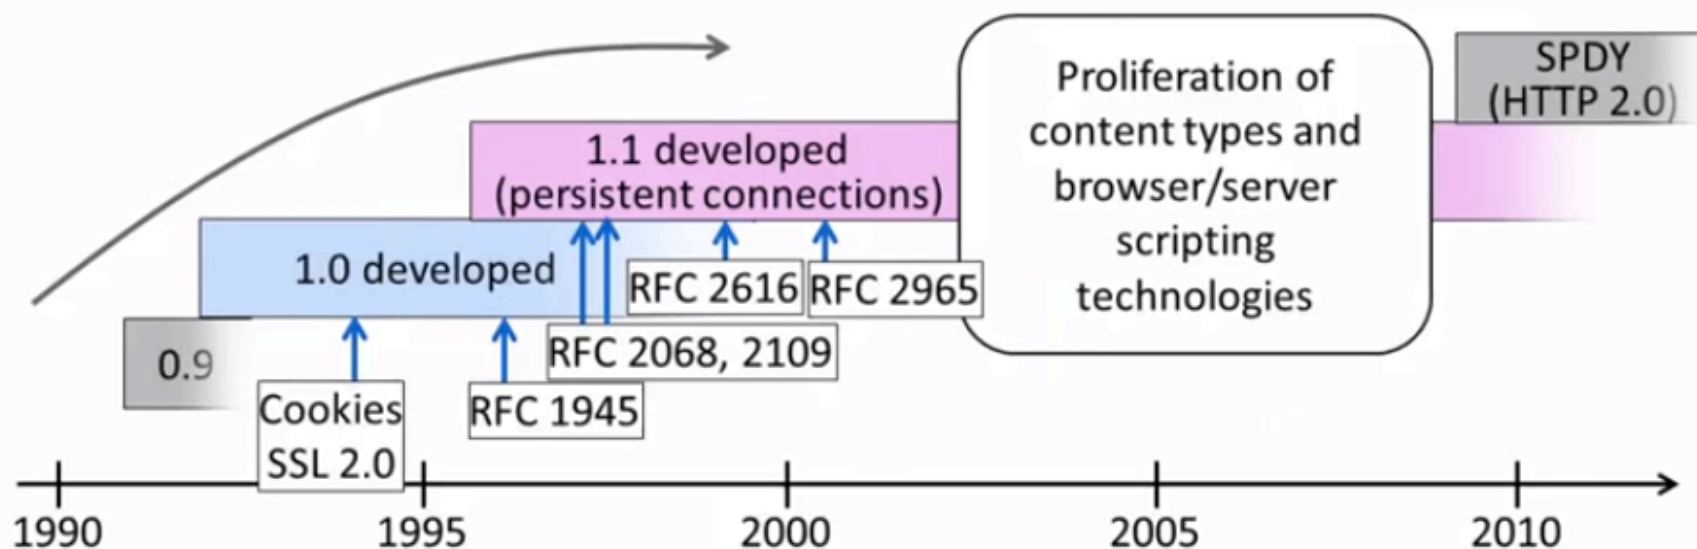
Static vs Dynamic Web pages

- Static web page is a file contents, e.g., image
- Dynamic web page is the result of program execution
 - Javascript on client, PHP on server, or both



Evolution of HTTP

- Consider security (SSL/TLS for HTTPS) later



HTTP Protocol

- Originally a simple protocol, with many options added over time
 - Text-based commands, headers
- Try it yourself:
 - As a “browser” fetching a URL
 - Run “telnet en.wikipedia.org 80”
 - Type “GET /wiki/Vegemite HTTP/1.0” to server followed by a blank line
 - Server will return HTTP response with the page contents (or other info)

HTTP Protocol

- Commands used in the request

	Method	Description
Fetch page →	GET	Read a Web page
	HEAD	Read a Web page's header
Upload data →	POST	Append to a Web page
	PUT	Store a Web page
	DELETE	Remove the Web page
	TRACE	Echo the incoming request
	CONNECT	Connect through a proxy
	OPTIONS	Query options for a page

HTTP Protocol

- Codes returned with the response

Code	Meaning	Examples
1xx	Information	100 = server agrees to handle client's request
2xx	Success	200 = request succeeded; 204 = no content present
3xx	Redirection	301 = page moved; 304 = cached page still valid
4xx	Client error	403 = forbidden page; 404 = page not found
5xx	Server error	500 = internal server error; 503 = try again later

Yes! →

HTTP Protocol

- Many header fields specify capabilities and content
 - Eg., Content-Type: text/html, Cookie:lect=9-1-http


Function	Example Headers
Browser capabilities (client → server)	User-Agent, Accept, Accept-Charset, Accept-Encoding, Accept-Language
Caching related (mixed directions)	If-Modified-Since, If-None-Match, Date, Last-Modified, Expires, Cache-Control, ETag
Browser context (client → server)	Cookie, Referer, Authorization, Host
Content delivery (server → client)	Content-Encoding, Content-Length, Content-Type, Content-Language, Content-Range, Set-Cookie

HTTP Performance


- **PLT (Page Load Time)**
 - PLT is the key measure of web performance
 - From click until user sees page
 - Small increases in PLT decrease sales
 - PLT depends on many factors
 - Structure of page/content
 - HTTP (and TCP!) protocol
 - Network RTT and bandwidth

Early Performance

- HTTP/1.0 uses one TCP connection to fetch one web resource
 - Made HTTP very easy to build
 - But gave fairly poor PLT ...




Early Performance



Early Performance

- Many reasons why PLT is larger than necessary
 - Sequential request/responses, even when to different servers
 - Multiple TCP connection setups to the same server
 - Multiple TCP slow-start phases
- Network is not used effectively
 - Worse with many small resources / page



Ways to Decrease PLT

1. Reduce content size for transfer
 - Smaller images, gzip
 2. Change HTTP to make better use of available bandwidth
 3. Change HTTP to avoid repeated transfers of the same content
 - Caching, and proxies
 4. Move content closer to client
 - CDNs [later]
-
- This time
- Next time
- Later


Parallel Connections

- One simple way to reduce PLT
 - Browser runs multiple (8, say) HTTP instances in parallel
 - Server is unchanged; already handled concurrent requests for many clients
- How does this help?
 - Single HTTP wasn't using network much ...
 - So parallel connections aren't slowed much
 - Pulls in completion time of last fetch

Persistent Connections

- Parallel connections compete with each other for network resources
 - 1 parallel client \approx 8 sequential clients?
 - Exacerbates network bursts, and loss
- Persistent connection alternative
 - Make 1 TCP connection to 1 server
 - Use it for multiple HTTP requests

Comparison



Overview of Persistent Connections

- Widely used as part of HTTP/1.1
 - Supports optional pipelining
 - PLT benefits depending on page structure, but easy on network
- Issues with persistent connections
 - How long to keep TCP connection?
 - Can it be slower? (Yes. But why?)

Next Class

- An overview of the HTTP formats
- An overview of Catching and Proxies